

LITCOIN NPL CHALLENGE

Resultados Iniciales del LitCoin NPL Challenge



PROBLEMA CIENTIFICO

El reto consiste en analizar los títulos y resúmenes de artículos biomédicos para identificar entidades como genes, enfermedades o productos químicos, y predecir las relaciones entre esas entidades, como asociaciones, correlaciones o interacciones. Los participantes deben entrenar modelos de procesamiento de lenguaje natural (NLP) para realizar esta tarea de forma automática, utilizando los datos proporcionados en formato CSV para entrenar y evaluar sus modelos.



Resultados Iniciales

OBJETIVOS

- + Desarrollar un modelo que prediga con precisión las relaciones entre entidades biomédicas en resúmenes de artículos científicos, como asociaciones, correlaciones o interacciones.
- + Evaluar y optimizar el rendimiento del modelo para garantizar que sea eficiente y escalable, permitiendo analizar grandes volúmenes de datos biomédicos.



LOS MODELOS A UTILIZAR

LSTM

RNN

+

Transformers

+

Por que LTSM


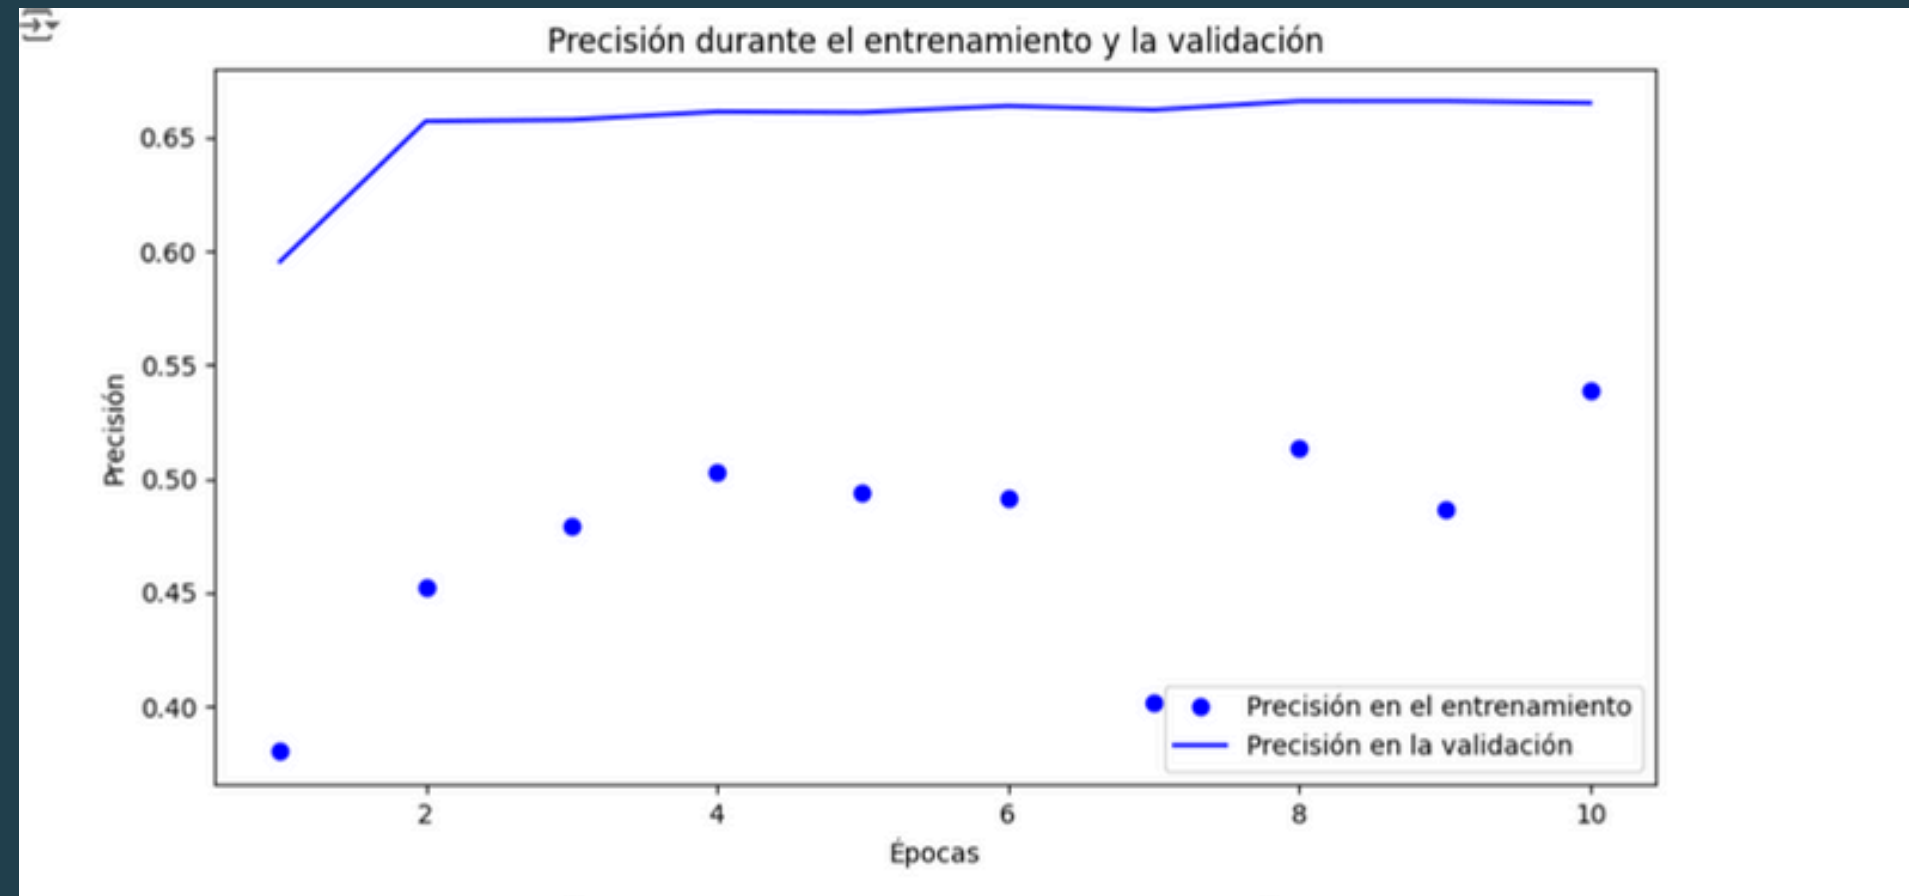
- + Captura de dependencias a largo plazo
- + Manejo efectivo de secuencias de texto
- + Quitar valores nulos



Resultados de LTSM

Precisión Entrenamiento 66%

Precisión Validacion 55%



Parametros

LSTM

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	?	0 (unbuilt)
bidirectional_1 (Bidirectional)	?	0 (unbuilt)
lstm_4 (LSTM)	?	0 (unbuilt)
dropout_2 (Dropout)	?	0 (unbuilt)
dense_2 (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)
Trainable params: 0 (0.00 B)
Non-trainable params: 0 (0.00 B)

Epoch 1/10
1848/1848 — 821s 439ms/step - accuracy: 0.5263 - loss: 1.3512 - val_accuracy: 0.3802 - val_loss: 1.4311

Epoch 2/10
1848/1848 — 855s 436ms/step - accuracy: 0.6571 - loss: 0.9120 - val_accuracy: 0.4521 - val_loss: 1.3434

Epoch 3/10
1848/1848 — 861s 435ms/step - accuracy: 0.6556 - loss: 0.8655 - val_accuracy: 0.4793 - val_loss: 1.4065

Epoch 4/10
1848/1848 — 891s 451ms/step - accuracy: 0.6593 - loss: 0.8360 - val_accuracy: 0.5026 - val_loss: 1.3387

Epoch 5/10
1848/1848 — 834s 436ms/step - accuracy: 0.6608 - loss: 0.8199 - val_accuracy: 0.4939 - val_loss: 1.3217

Epoch 6/10
1848/1848 — 865s 438ms/step - accuracy: 0.6648 - loss: 0.8059 - val_accuracy: 0.4916 - val_loss: 1.3124

Epoch 7/10
1848/1848 — 863s 438ms/step - accuracy: 0.6642 - loss: 0.7999 - val_accuracy: 0.4018 - val_loss: 1.2907

Epoch 8/10
1848/1848 — 860s 438ms/step - accuracy: 0.6653 - loss: 0.7979 - val_accuracy: 0.5137 - val_loss: 1.3434

Epoch 9/10
1848/1848 — 861s 437ms/step - accuracy: 0.6640 - loss: 0.7978 - val_accuracy: 0.4869 - val_loss: 1.3016

Epoch 10/10



Por qué **RNN**



Procesamiento secuencial de texto



Captura de dependencias contextuales locales




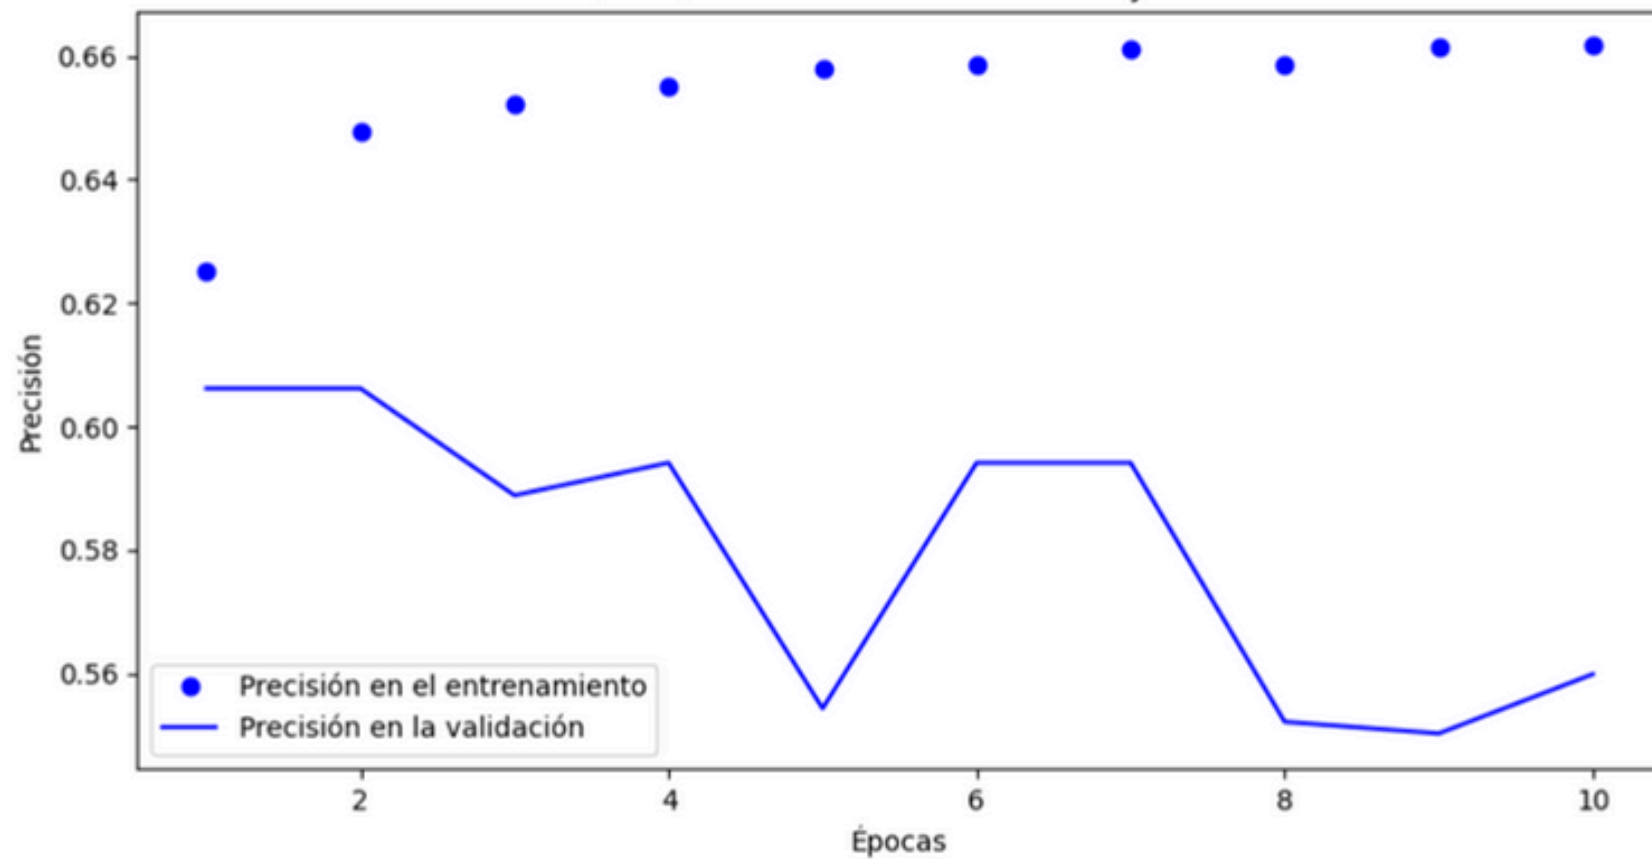
Resultados de RNN



Precisión Entrenamiento: %66

Precisión Validacion 61%

Precisión (RNN) durante el entrenamiento y la validación



RNN

bidirectional (Bidirectional)	?	0 (unbuilt)
simple_rnn_1 (SimpleRNN)	?	0 (unbuilt)
dropout (Dropout)	?	0 (unbuilt)
dense (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

1848/1848 ————— 298s 157ms/step - accuracy: 0.5731 - loss: 1.2182 - val_accuracy: 0.4835 - val_loss: 1.3446

Epoch 2/10

1848/1848 ————— 286s 155ms/step - accuracy: 0.6438 - loss: 0.8976 - val_accuracy: 0.4156 - val_loss: 1.2631

Epoch 3/10

1848/1848 ————— 324s 156ms/step - accuracy: 0.6541 - loss: 0.8499 - val_accuracy: 0.4578 - val_loss: 1.2461

Epoch 4/10

1848/1848 ————— 286s 155ms/step - accuracy: 0.6557 - loss: 0.8270 - val_accuracy: 0.4484 - val_loss: 1.2091

Epoch 5/10

1848/1848 ————— 287s 155ms/step - accuracy: 0.6588 - loss: 0.8194 - val_accuracy: 0.4277 - val_loss: 1.2207

Epoch 6/10

1848/1848 ————— 320s 155ms/step - accuracy: 0.6579 - loss: 0.8100 - val_accuracy: 0.4547 - val_loss: 1.2114

Epoch 7/10

1848/1848 ————— 321s 154ms/step - accuracy: 0.6598 - loss: 0.8058 - val_accuracy: 0.4330 - val_loss: 1.1945

Epoch 8/10

1848/1848 ————— 286s 155ms/step - accuracy: 0.6559 - loss: 0.8028 - val_accuracy: 0.4394 - val_loss: 1.2217

Epoch 9/10

1848/1848 ————— 286s 155ms/step - accuracy: 0.6637 - loss: 0.7938 - val_accuracy: 0.3866 - val_loss: 1.1973

Epoch 10/10

1848/1848 ————— 323s 155ms/step - accuracy: 0.6595 - loss: 0.7940 - val_accuracy: 0.4263 - val_loss: 1.1788

Por qué **TRANSFORMERS**



**Optimizado para datos
secuenciales**



**Mecanismos de atención para
enfocar en las partes más
relevantes**





Resultados de TRANSFORMERS



Precisión Entrenamiento: %78

Precisión Validación: 79%



Valores de entrenamiento

TRANSFORMERS

100% 1000/1000 20:29-20:30 1.2894s

```
("loss": 1.8087, "grad_norm": 5.882642911151244, "learning_rate": 1.7623889238892389e-05, "epoch": 0.45)
("loss": 0.8171, "grad_norm": 13.21344537344171, "learning_rate": 1.4943298432984329e-05, "epoch": 0.89)
```

```
llama_model_load_weights(local_filepath=Path('huggingface/transformers/pytorch_model.bin'), backend=BackendClassification, num_gpus=undefined) Precision is ill-defined and being set to 0.0 in labels with no predicted samples. use 'zero_division' parameter to control this behavior.
  _print_perf(average, modifier, f'({metric.capitalize()}) is', lambda result:
("eval_loss": 0.708680605946793, "eval_accuracy": 0.737882768888788, "eval_f1": 0.738878299643775, "eval_precision": 0.7246888413104342, "eval_recall": 0.737882768888788, "eval_runtime": 35.4181, "eval_samples_per_second": 23.407, "eval_steps_per_second": 2.429, "epoch": 1.4)
("loss": 0.7981, "grad_norm": 12.565623854492188, "learning_rate": 1.3876428764287642e-05, "epoch": 1.94)
("loss": 0.4687, "grad_norm": 11.485288135412124, "learning_rate": 0.4952388923889239e-05, "epoch": 1.79)
```

```
llama_model_load_weights(local_filepath=Path('huggingface/transformers/pytorch_model.bin'), backend=BackendClassification, num_gpus=undefined) Precision is ill-defined and being set to 0.0 in labels with no predicted samples. use 'zero_division' parameter to control this behavior.
  _print_perf(average, modifier, f'({metric.capitalize()}) is', lambda result:
("eval_loss": 0.4477281835421872, "eval_accuracy": 0.76887881811288, "eval_f1": 0.76887879389388, "eval_precision": 0.768111111127829, "eval_recall": 0.76887881811288, "eval_runtime": 35.1994, "eval_samples_per_second": 23.421, "eval_steps_per_second": 2.429, "epoch": 1.4)
("loss": 0.5684, "grad_norm": 11.84858813444814, "learning_rate": 5.119847638476384e-06, "epoch": 2.23)
("loss": 0.4079, "grad_norm": 9.74979498807372, "learning_rate": 2.1428571428571428e-06, "epoch": 2.48)
```

```
llama_model_load_weights(local_filepath=Path('huggingface/transformers/pytorch_model.bin'), backend=BackendClassification, num_gpus=undefined) Precision is ill-defined and being set to 0.0 in labels with no predicted samples. use 'zero_division' parameter to control this behavior.
  _print_perf(average, modifier, f'({metric.capitalize()}) is', lambda result:
("eval_loss": 0.426473685488547, "eval_accuracy": 0.76888483998154, "eval_f1": 0.76888488816751, "eval_precision": 0.768588854393127, "eval_recall": 0.76888483998154, "eval_runtime": 30.8154, "eval_samples_per_second": 24.33, "eval_steps_per_second": 2.443, "epoch": 3.4)
("train_runtime": 3579.8134, "train_samples_per_second": 7.585, "train_steps_per_second": 0.878, "train_loss": 0.4915263579268857, "epoch": 3.4)

train_output(global_step=1048, training_loss=0.4915263579268857, metrics={"train_runtime": 3579.8134, "train_samples_per_second": 7.585, "train_steps_per_second": 0.878, "total_flos": 2768964811193440.0, "train_loss": 0.4915263579268857, "epoch": 3.4})
```


REFERENCIAS

1. Lee, J., Yoon, W., Kim, S., Kim, D., Kim, S., So, C. H., & Kang, J. (2020). BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4), 1234–1240. Disponible en <https://academic.oup.com/bioinformatics/article/36/4/1234/5566506?login=false>.
2. Muchene, L., & Safari, W. (2021). Two-stage topic modelling of scientific publications: A case study of University of Nairobi, Kenya. *PLoS One*, 16(1), e0243208. Disponible en <https://doi.org/10.1371/journal.pone.0243208>.
3. James, H. (2022). RNNs and LSTMs. Disponible en <https://web.stanford.edu/~jurafsky/slp3/9.pdf>.

 Analisis Exploratorio



ANALISIS DE LOS ABSTRACTS




```
Resumen de variables:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 3 columns):
#      Column          Non-Null Count  Dtype
---  -
0     abstract_id    400 non-null    int64
1     title          400 non-null    object
2     abstract       400 non-null    object
dtypes: int64(1), object(2)
```





Analisis Exploratorio

Histograma de Longitud de Resúmenes




Histograma de Longitud de Títulos



 Analisis Exploratorio




ANALISIS DE LAS ENTIDADES




```
Resumen de variables:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13636 entries, 0 to 13635
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id              13636 non-null  int64
1   abstract_id     13636 non-null  int64
2   offset_start    13636 non-null  int64
3   offset_finish   13636 non-null  int64
4   type            13636 non-null  object
5   mention         13636 non-null  object
6   entity_ids      13636 non-null  object
dtypes: int64(4), object(3)
memory usage: 745.8+ KB
None
```

Ajustes en los datos



Distribución de Offset Start por Tipo de Entidad




 Analisis Exploratorio




ANALISIS DE LAS RELACIONES



Distribución de Novedad de Relaciones en relations_train




Distribución de Tipos de Relaciones en relations_train



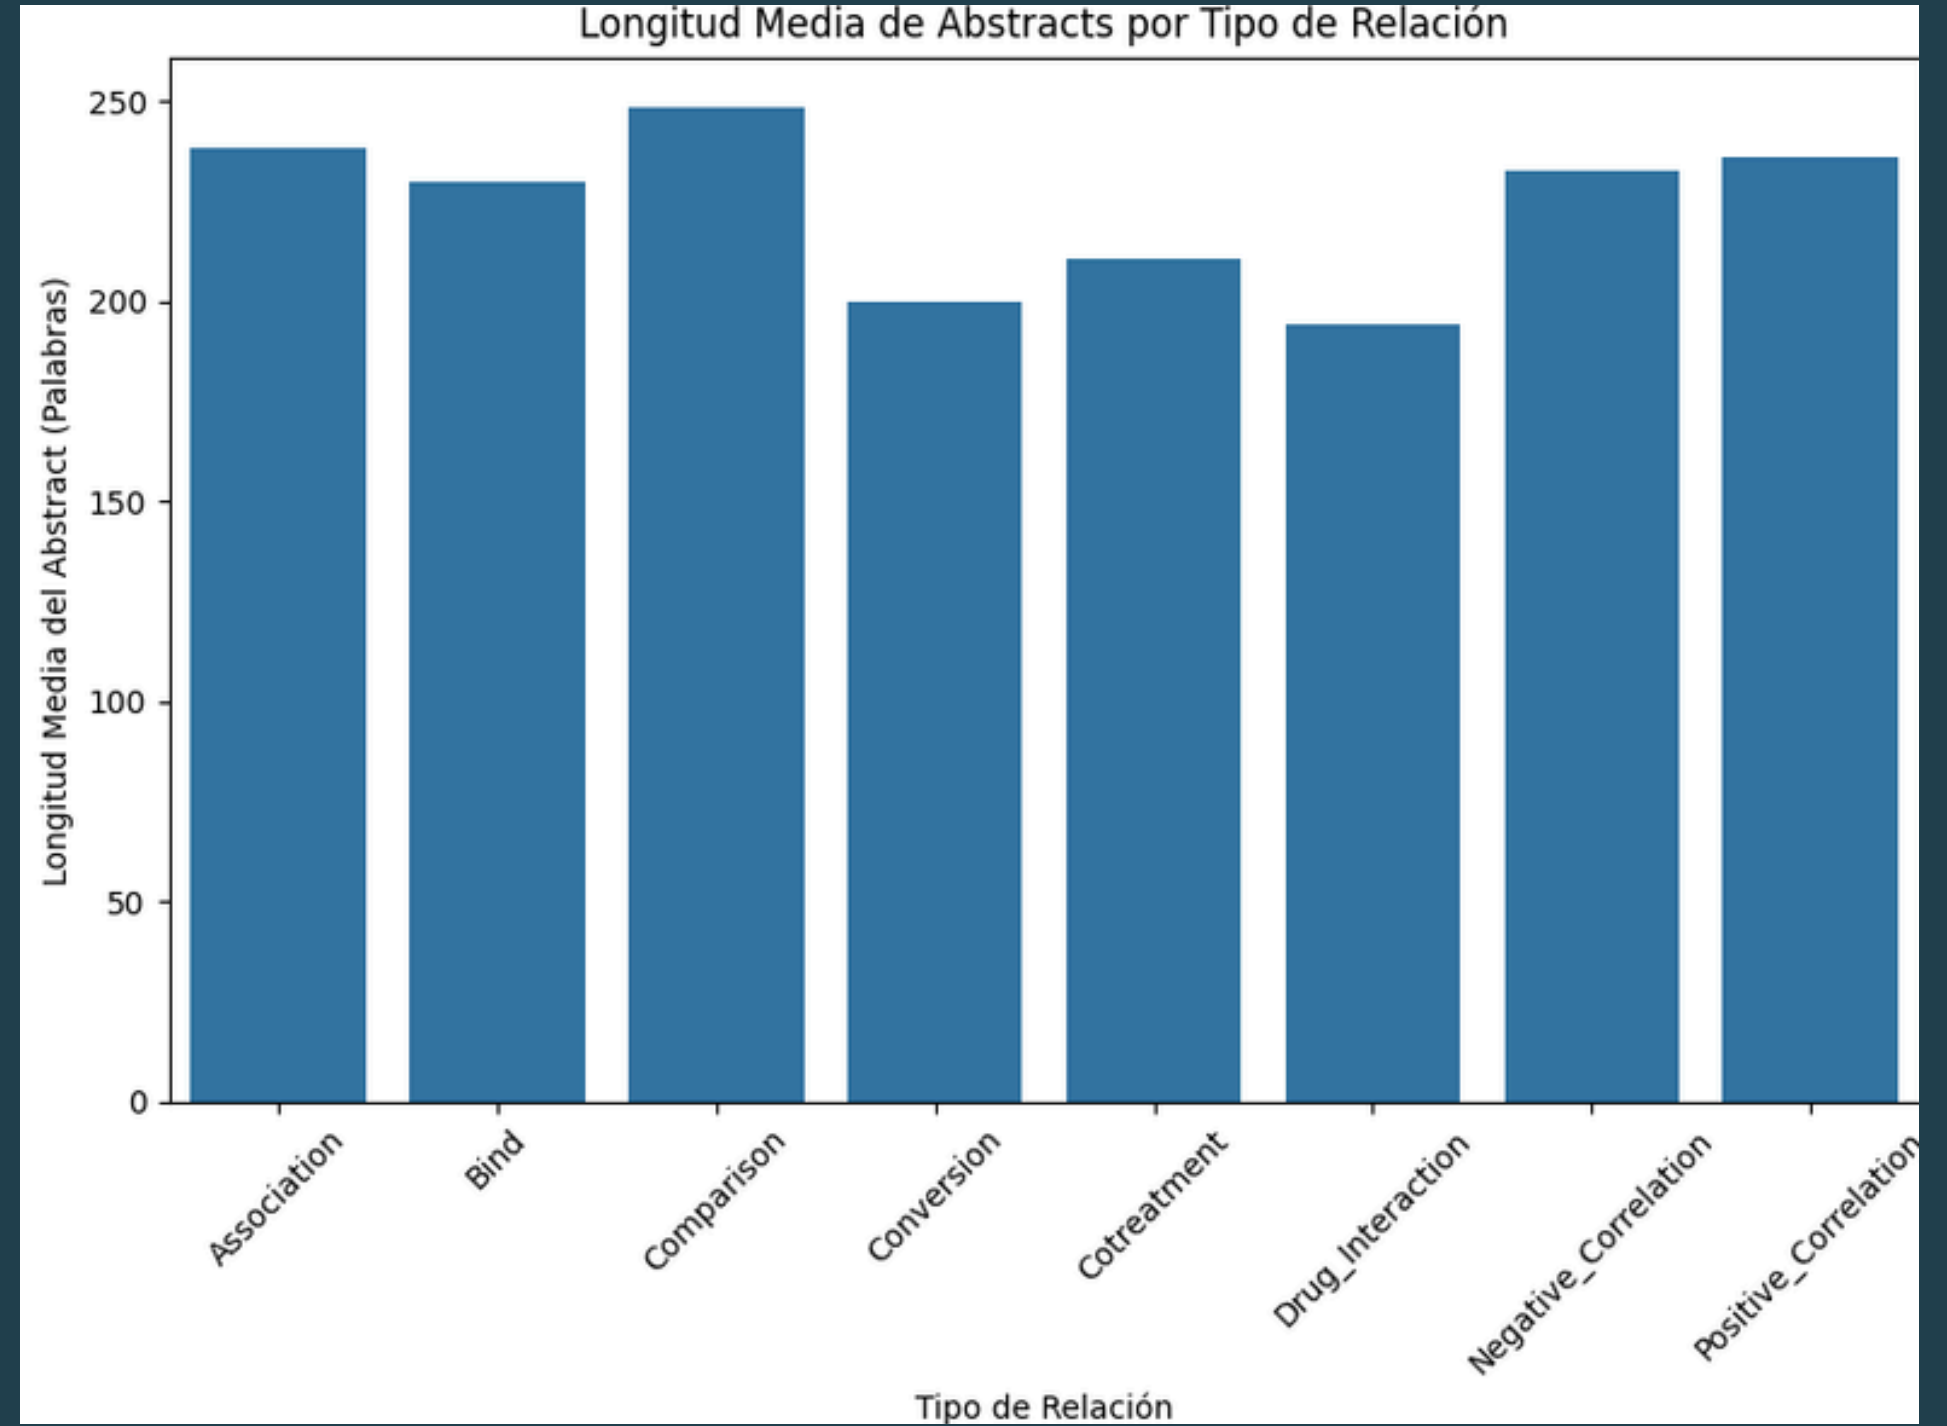
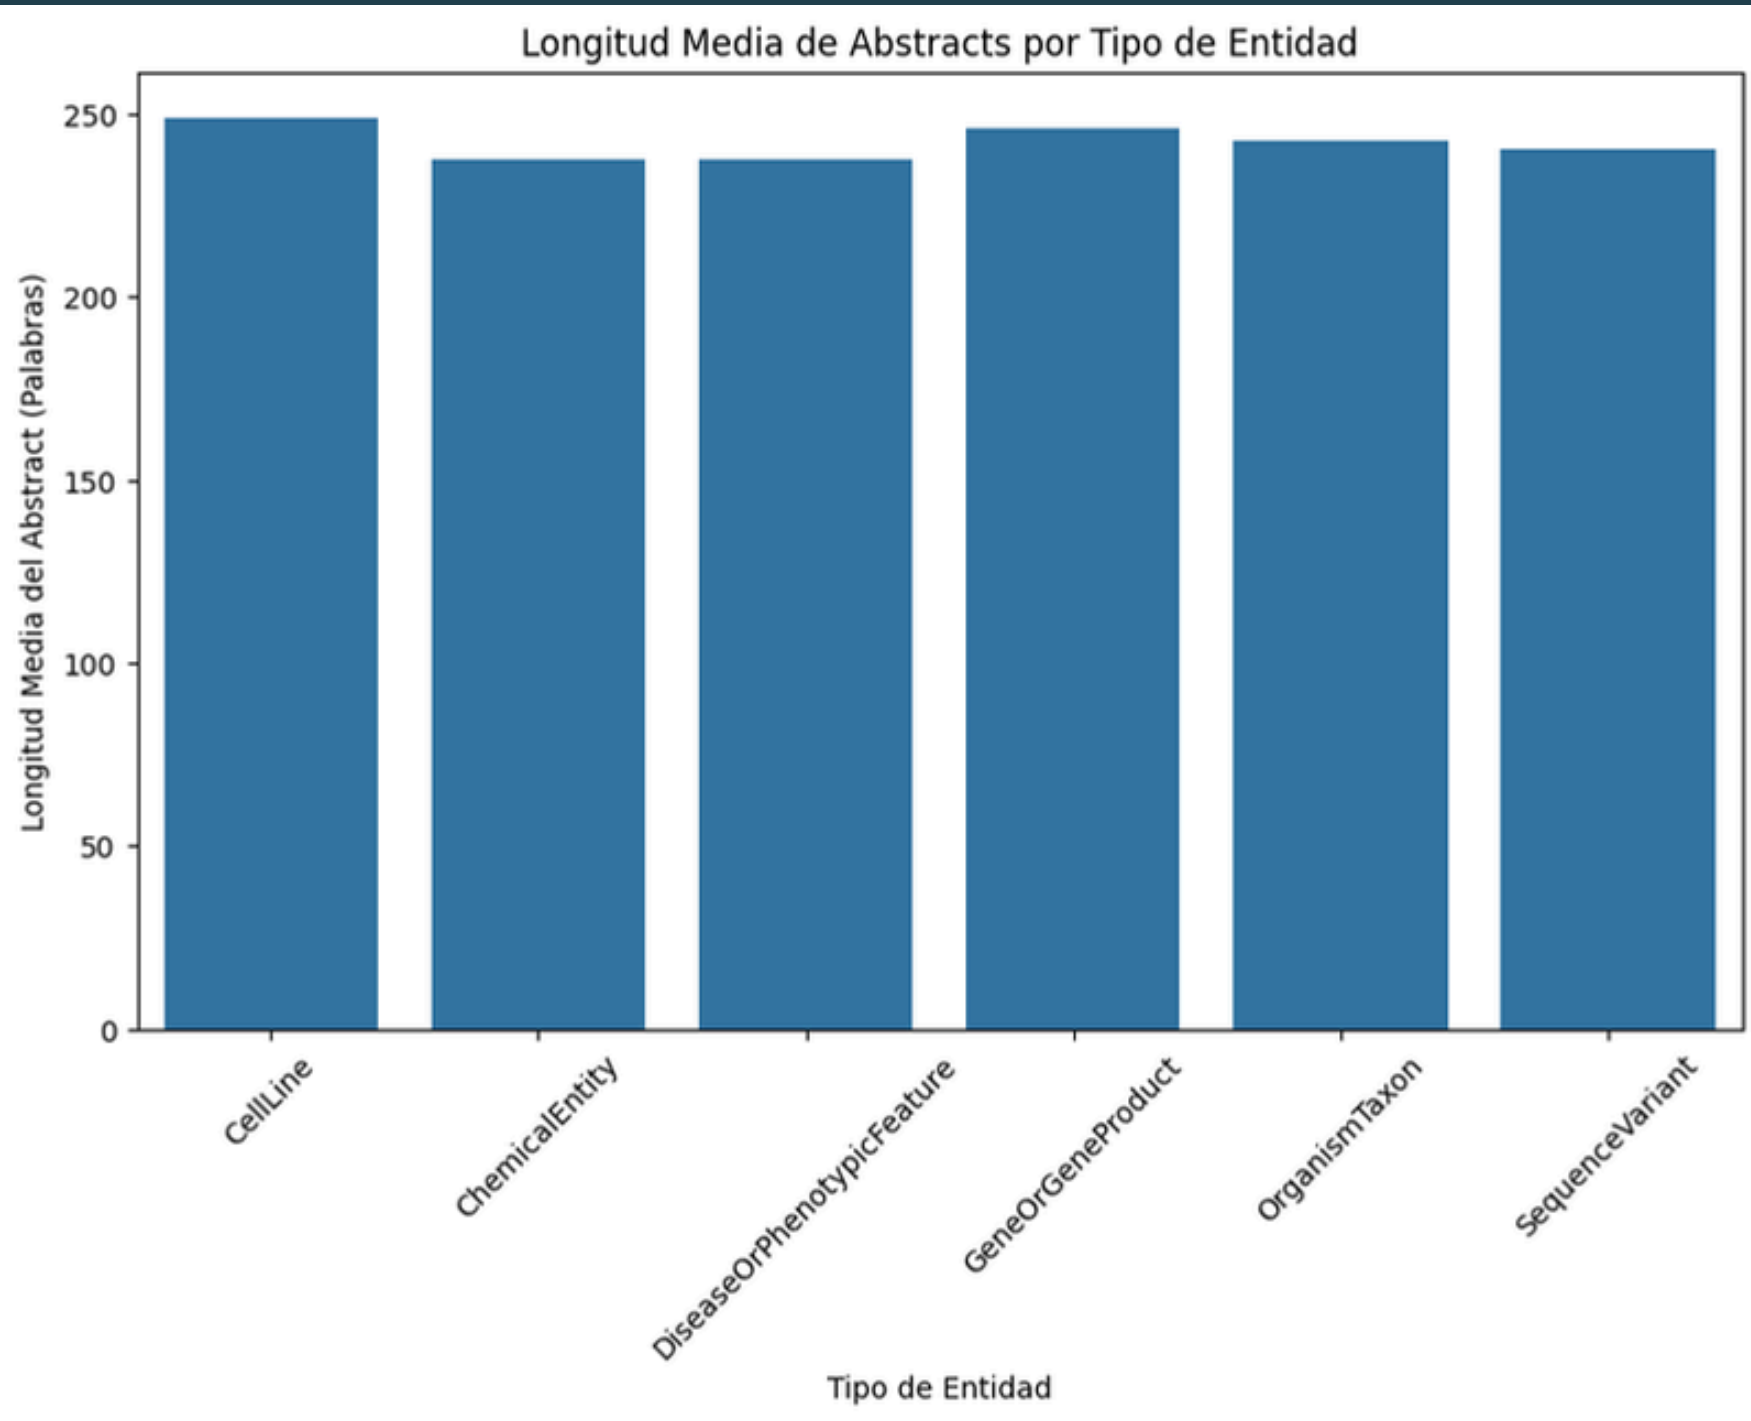
 Analisis Exploratorio




ANALISIS DE MEZCLADO



```
Resumen de variables:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4280 entries, 0 to 4279
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               4280 non-null   int64
1   abstract_id      4280 non-null   int64
2   type             4280 non-null   object
3   entity_1_id      4280 non-null   object
4   entity_2_id      4280 non-null   object
5   novel            4280 non-null   object
dtypes: int64(2), object(4)
memory usage: 200.8+ KB
None
```



Distribución de Tipos de Entidades por Tipo de Relación





Resultados iniciales

GRACIAS +

